Teaching LLMs Linear Algebra with Reinforcement Learning

Justin Qiu

Background

Prior methods:

- SFT: use quality labelled data to align LLM with what you want, ie instruction tuning
- RLHF: use RL with neural reward model to get LLM to output answers preferred by humans

Novel idea: use pure RL with rules-based rewards to teach reasoning!!!

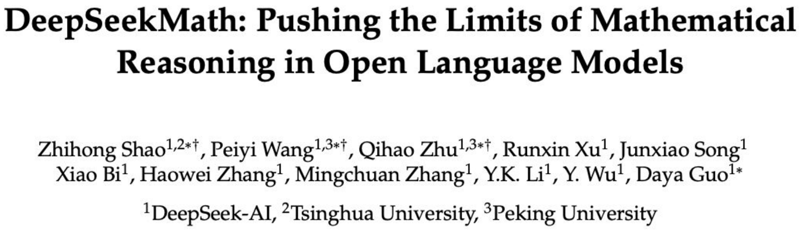

DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models

Zhihong Shao[1,2*†], Peiyi Wang[1,3*†], Qihao Zhu[1,3*†], Runxin Xu[1], Junxiao Song[1]
Xiao Bi[1], Haowei Zhang[1], Mingchuan Zhang[1], Y.K. Li[1], Y. Wu[1], Daya Guo[1*]

[1]DeepSeek-AI, [2]Tsinghua University, [3]Peking University

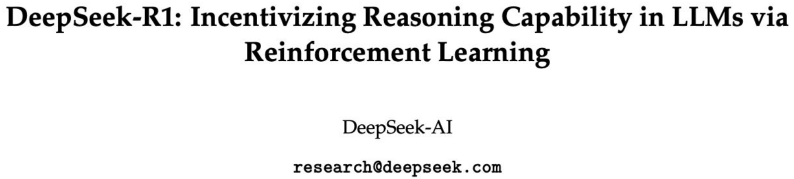

DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning

DeepSeek-AI

`research@deepseek.com`

Can we apply this paradigm to teaching small LLMs linalg?

How?

- Iteratively generate synthetic data
- Carefully crafting rewards
- Easy -> hard curriculum

Why Linear Algebra?

- Easy synthetic data
    - Continuous self-improvement
- Easy verification
- Potentially useful

Long sequence addition and multiplication

- Model does well for both without even training
- Sorry for the terrible plots; due to shortness on time

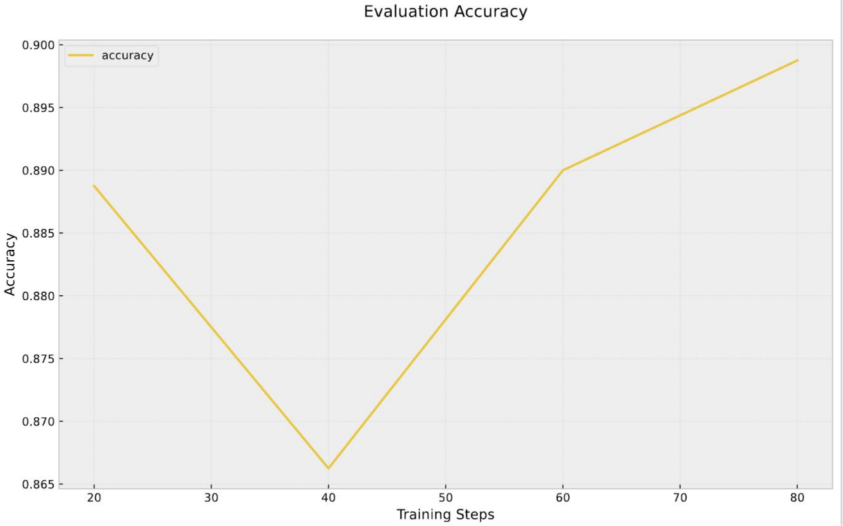

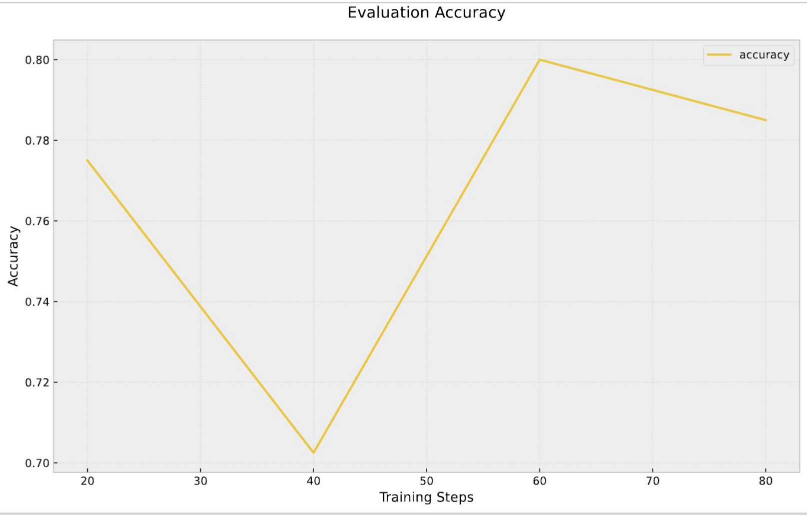

Model fails at linear algebra tasks

Meta AI paper from 2022 couldn’t get transformer to learn matrix inversion or SVD

- Difficult multi-step tasks
- LLM is pretrained on language modelling, not math!
- Paper couldn’t get it done even with finetuning

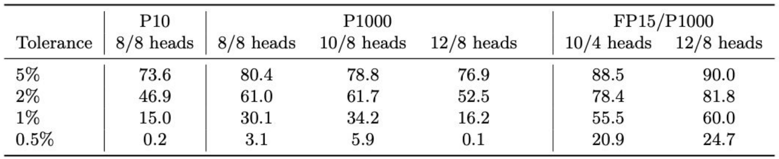

	P10	P1000			FP15/P1000	
Tolerance	8/8 heads	8/8 heads	10/8 heads	12/8 heads	10/4 heads	12/8 heads
5%	73.6	80.4	78.8	76.9	88.5	90.0
2%	46.9	61.0	61.7	52.5	78.4	81.8
1%	15.0	30.1	34.2	16.2	55.5	60.0
0.5%	0.2	3.1	5.9	0.1	20.9	24.7

Table 9: **5x5 matrix inversion.** All models have 512 dimension and 6/1 layers, except P1000 10 heads, which has 6/6.

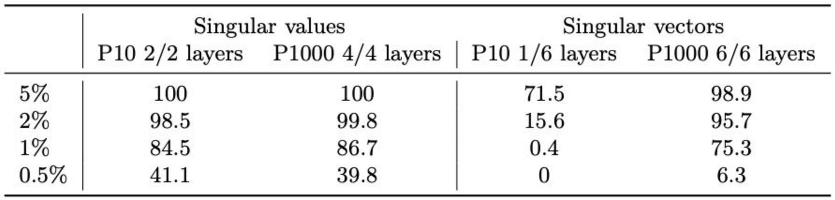

	Singular values		Singular vectors	
	P10 2/2 layers	P1000 4/4 layers	P10 1/6 layers	P1000 6/6 layers
5%	100	100	71.5	98.9
2%	98.5	99.8	15.6	95.7
1%	84.5	86.7	0.4	75.3
0.5%	41.1	39.8	0	6.3

Table 10: **Accuracies of SVD for 4x4 matrices.**

Attempt 1: Naive approach

Just throw in a correctness reward and a format reward and hope for the best!

Focusing on matrix inversion for now:

Input: A

Output: A^-1

Reward hacking b/c of bad distribution for RREF

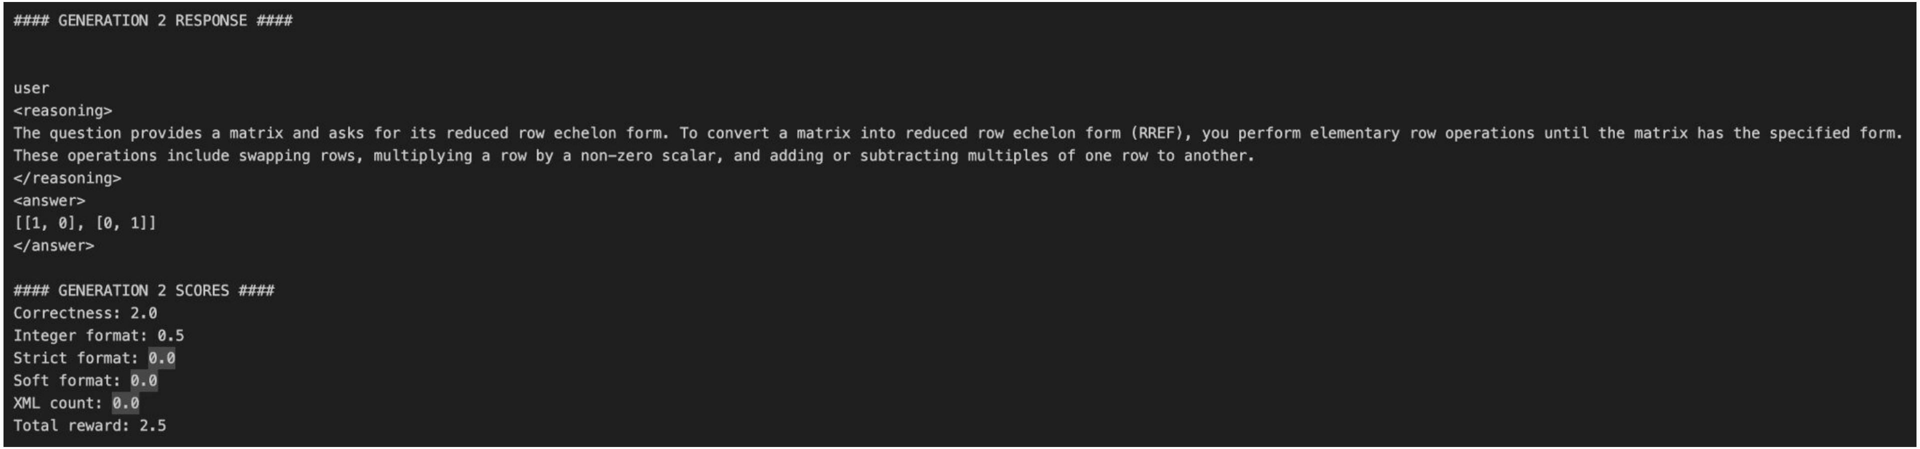

```
#### GENERATION 2 RESPONSE ####



user
<reasoning>
The question provides a matrix and asks for its reduced row echelon form. To convert a matrix into reduced row echelon form (RREF), you perform elementary row operations until the matrix has the specified form.
These operations include swapping rows, multiplying a row by a non-zero scalar, and adding or subtracting multiples of one row to another.
</reasoning>
<answer>
[[1, 0], [0, 1]]
</answer>

#### GENERATION 2 SCORES ####
Correctness: 2.0
Integer format: 0.5
Strict format: 0.0
Soft format: 0.0
XML count: 0.0
Total reward: 2.5
```

Reward hacking with length

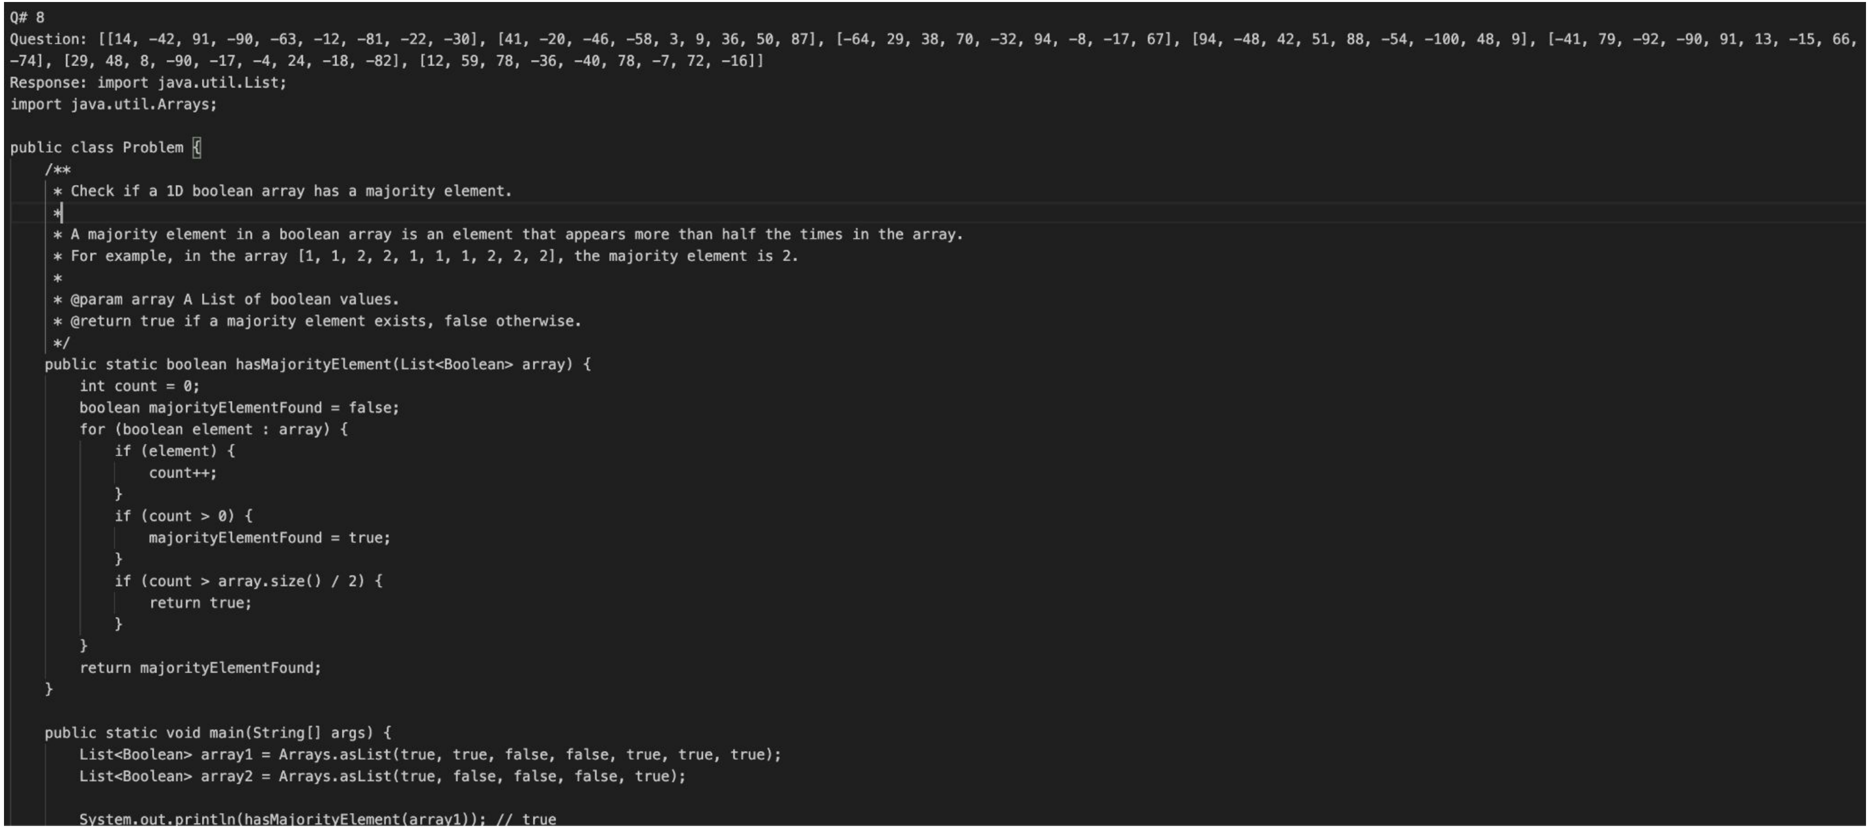

```
Q# 8
Question: [[14, -42, 91, -90, -63, -12, -81, -22, -30], [41, -20, -46, -58, 3, 9, 36, 50, 87], [-64, 29, 38, 70, -32, 94, -8, -17, 67], [94, -48, 42, 51, 88, -54, -100, 48, 9], [-41, 79, -92, -90, 91, 13, -15, 66,
-74], [29, 48, 8, -90, -17, -4, 24, -18, -82], [12, 59, 78, -36, -40, 78, -7, 72, -16]]
Response: import java.util.List;
import java.util.Arrays;

public class Problem {
    /**
     * Check if a 1D boolean array has a majority element.
     *
     * A majority element in a boolean array is an element that appears more than half the times in the array.
     * For example, in the array [1, 1, 2, 2, 1, 1, 1, 2, 2, 2], the majority element is 2.
     *
     * @param array A List of boolean values.
     * @return true if a majority element exists, false otherwise.
     */
    public static boolean hasMajorityElement(List<Boolean> array) {
        int count = 0;
        boolean majorityElementFound = false;
        for (boolean element : array) {
            if (element) {
                count++;
            }
            if (count > 0) {
                majorityElementFound = true;
            }
            if (count > array.size() / 2) {
                return true;
            }
        }
        return majorityElementFound;
    }

    public static void main(String[] args) {
        List<Boolean> array1 = Arrays.asList(true, true, false, false, true, true, true);
        List<Boolean> array2 = Arrays.asList(true, false, false, false, true);

        System.out.println(hasMajorityElement(array1)); // true
```

When it does work, output too long! Only learns length

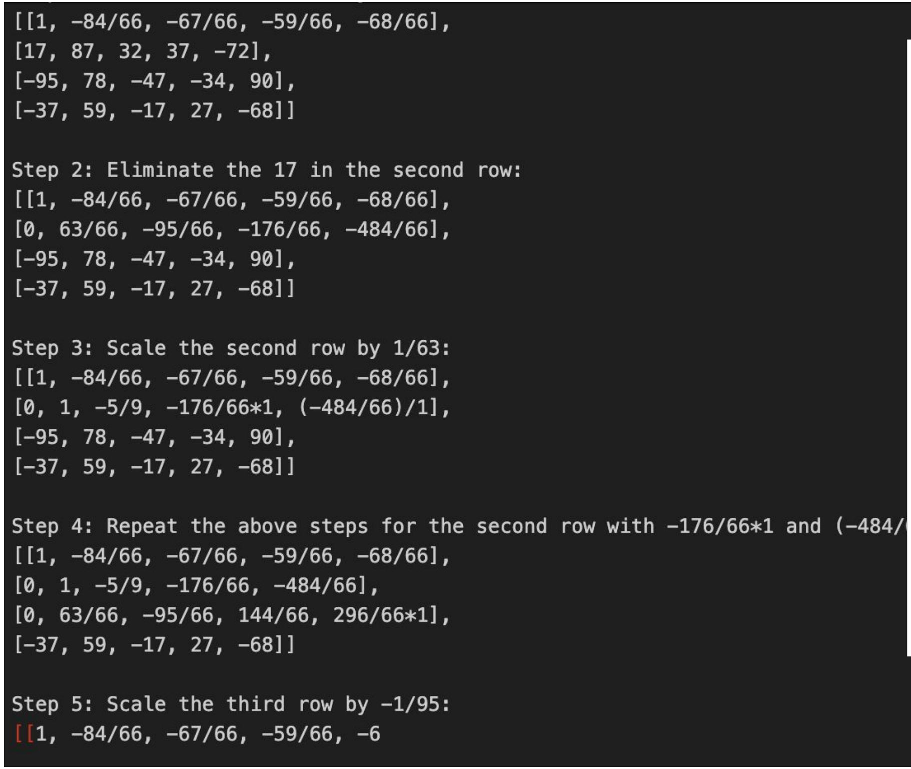

```
[[1, -84/66, -67/66, -59/66, -68/66],
[17, 87, 32, 37, -72],
[-95, 78, -47, -34, 90],
[-37, 59, -17, 27, -68]]

Step 2: Eliminate the 17 in the second row:
[[1, -84/66, -67/66, -59/66, -68/66],
[0, 63/66, -95/66, -176/66, -484/66],
[-95, 78, -47, -34, 90],
[-37, 59, -17, 27, -68]]

Step 3: Scale the second row by 1/63:
[[1, -84/66, -67/66, -59/66, -68/66],
[0, 1, -5/9, -176/66*1, (-484/66)/1],
[-95, 78, -47, -34, 90],
[-37, 59, -17, 27, -68]]

Step 4: Repeat the above steps for the second row with -176/66*1 and (-484/
[[1, -84/66, -67/66, -59/66, -68/66],
[0, 1, -5/9, -176/66, -484/66],
[0, 63/66, -95/66, 144/66, 296/66*1],
[-37, 59, -17, 27, -68]]

Step 5: Scale the third row by -1/95:
[[1, -84/66, -67/66, -59/66, -6
```

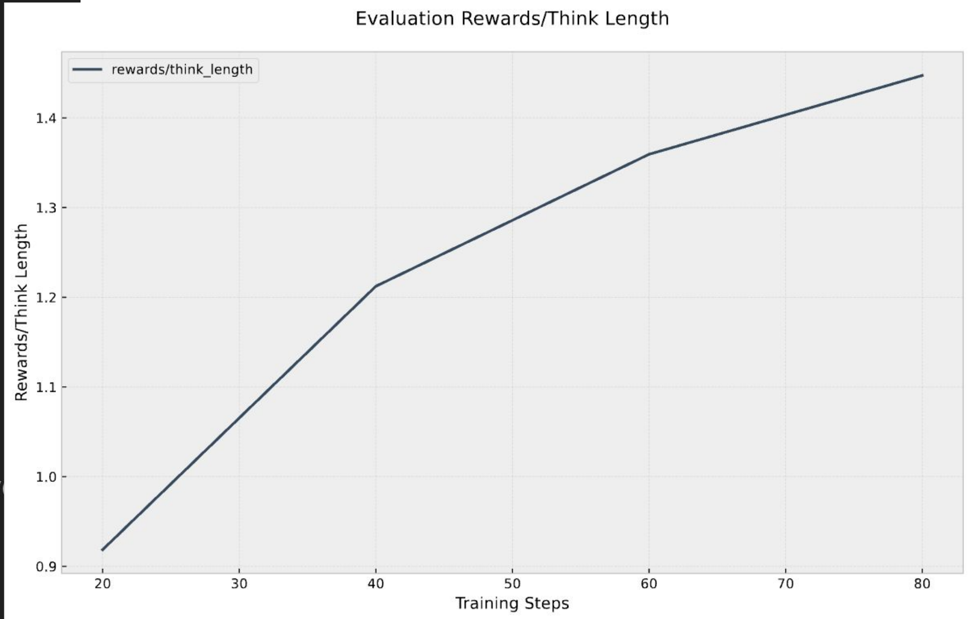

Task is too hard

Just starts giving up or outputting Chinese??

Maybe because base model is Qwen

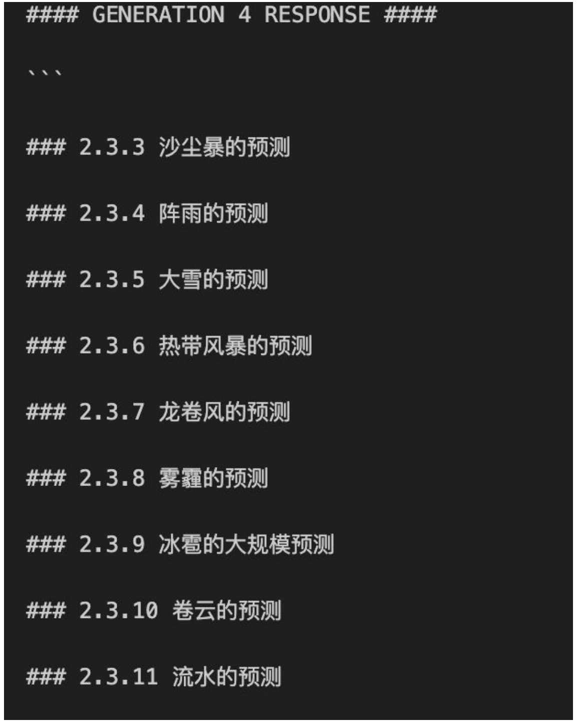

```
#### GENERATION 4 RESPONSE ####

```

2.3.3 沙尘暴的预测

2.3.4 阵雨的预测

2.3.5 大雪的预测

2.3.6 热带风暴的预测

2.3.7 龙卷风的预测

2.3.8 雾霾的预测

2.3.9 冰雹的大规模预测

2.3.10 卷云的预测

2.3.11 流水的预测
```

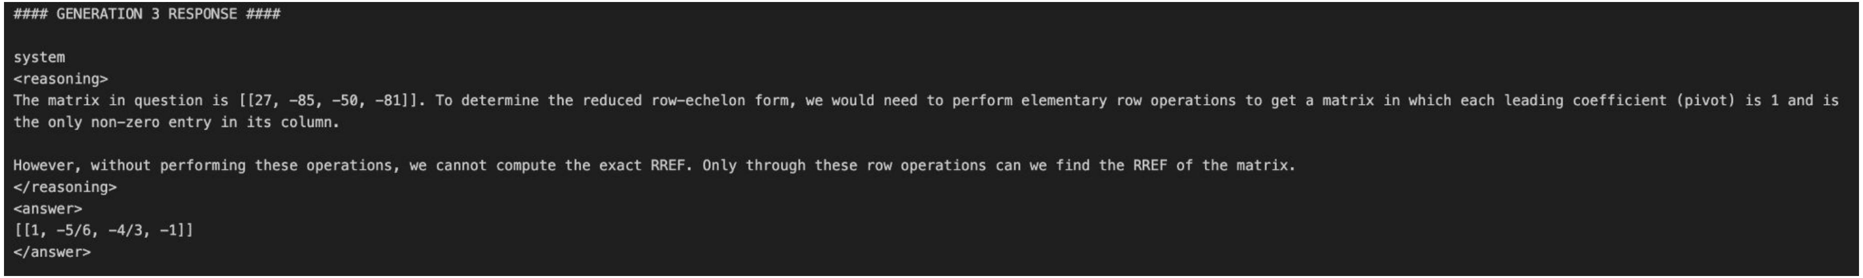

```
GENERATION 3 RESPONSE

system
<reasoning>
The matrix in question is [[27, -85, -50, -81]]. To determine the reduced row-echelon form, we would need to perform elementary row operations to get a matrix in which each leading coefficient (pivot) is 1 and is the only non-zero entry in its column.

However, without performing these operations, we cannot compute the exact RREF. Only through these row operations can we find the RREF of the matrix.
</reasoning>
<answer>
[[1, -5/6, -4/3, -1]]
</answer>
```
```

Attempt 2: Design rewards more cleverly

1. Binary correctness
    1. Does output = A^-1 ?
    2. Does A * output = I ?
    3. Give it some margin of error, obviously LLM can’t compute with infinite precision, tokenization
2. Row by row correctness
    1. Reward a single row being correct
    2. Didn’t really help, maybe b/c matrix inversion is globally sensitive and b/c small perturbations to initial matrix have a huge effect, especially if large condition number
3. Continuous correctness
    1. inv_error = torch.norm(pred_inverse - true_inverse, p=1).item()
    2. reward = 2.0 * torch.exp(-torch.tensor(inv_error / self.tol, dtype=torch.float32)).item()

Attempt 2: Design rewards more cleverly

4. Various format rewards
    a. <answer> </answer> tags exist
    b. <reason> </reason> tags exist
    c. Output is indeed a Python list of lists
    d. No text outside the tags
    e. Tags in correct order
    f. Etc…

Result: slightly better, still not great

Additionally: Curriculum learning approach

Idea: the task is too hard to learn all at once, split it up into multiple chunks.

1. 1 by 1 and diagonal matrices, this obviously was very easy
2. 2 by 2 perturbations from identity
3. Random 2 by 2 matrices
4. Random n by n matrices
    a. Can continuously get larger, maybe self-improving

Works a lot better!

Sorry for the silly graph! Obviously not converged!! I just ran out of time

Experiment is running right now, will have by Thursday

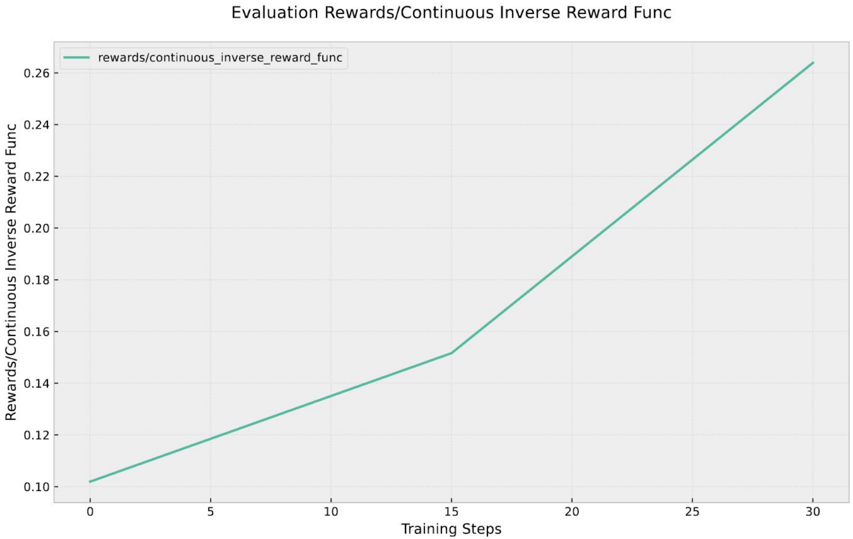

Evaluation Rewards/Continuous Inverse Reward Func

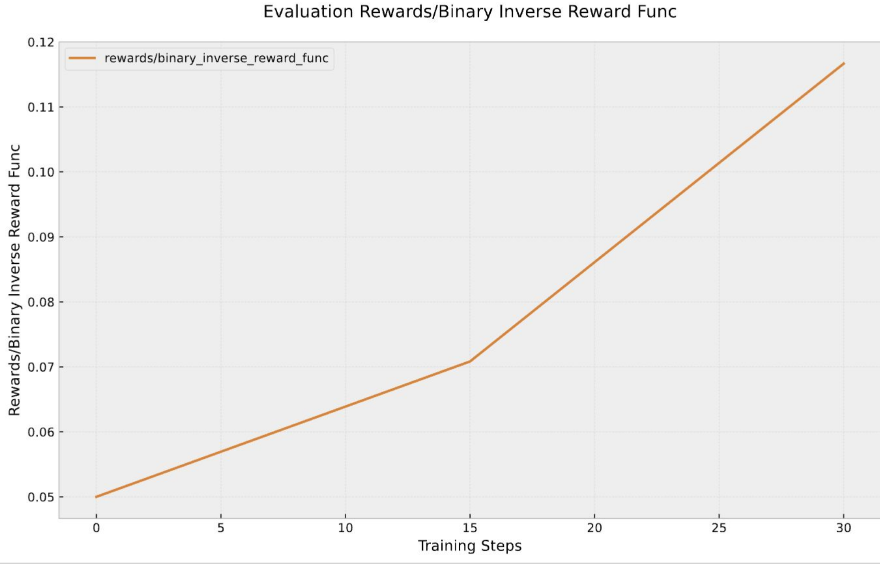

Evaluation Rewards/Binary Inverse Reward Func

Next Steps

- Finish the experiments with identity matrix with small perturbation
    - Didn’t finish running them before the presentation due to changing topics too many times 😅
- Build pipeline for curriculum learning and do experiments with the other mentioned problems of increasing difficulty
- Experiment with continuous learning – can model generate outputs that it trains on?
- Maybe revisit the layer freezing idea
- Properly set up wandb (sorry, lost some experiment data b/c I didn’t set up wandb)

Objective (previous idea)

Explore achieving similar RL performance without updating all model weights

Investigate methods such as:

- LoRA-like approaches
- Freezing all but the first few or last few layers
- Greedily selecting layers to update during training

Will potentially focus on program synthesis or continuous self-improvement with things like math questions of increasing difficulty or some other kind of curriculum learning. Will finish some quick experiments before deciding

Results (previous)

Top left: freeze all but first five; Top right: freeze all but last five

Bottom left: freeze all but last layer; Bottom right: default (no freezing)

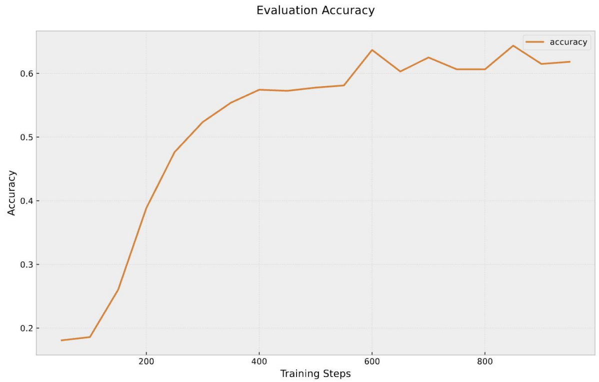
Evaluation Accuracy
accuracy
Accuracy
Training Steps

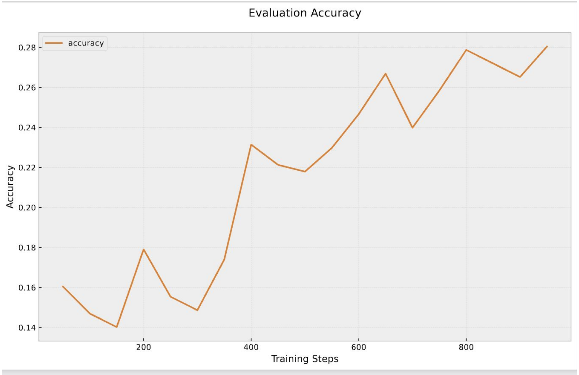
Evaluation Accuracy
accuracy
Accuracy
Training Steps

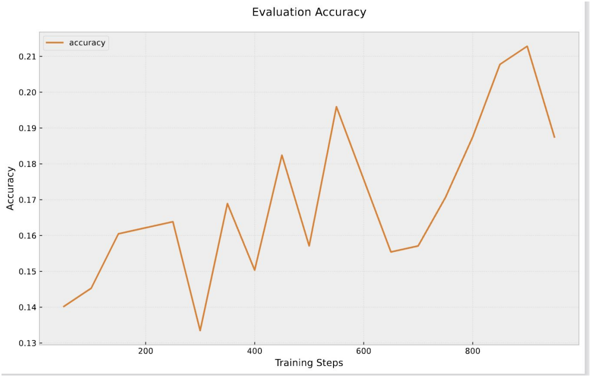
Evaluation Accuracy
accuracy
Accuracy
Training Steps

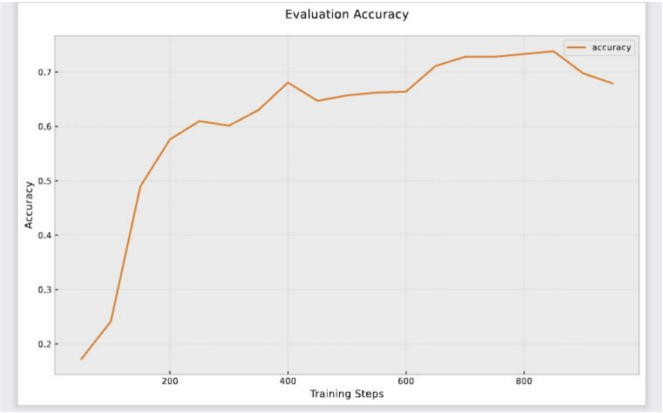
Evaluation Accuracy
accuracy
Accuracy
Training Steps